

```
In [1]: #PREDICTION OF SALES PRICE
import pandas as pd
```

```
In [2]: df=pd.read_csv('1. Regression - Module - (Housing Prices).csv')
```

```
In [3]: df['Condition of the House'].head(5)
```

Out[3]:

0	Fair
1	Fair
2	Fair
3	Excellent
4	Fair

Name: Condition of the House, dtype: object

```
In [4]: df['Sale Price'].describe()
```

Out[4]:

count	2.160900e+04
mean	5.401984e+05
std	3.673890e+05
min	7.500000e+04
25%	3.219500e+05
50%	4.500000e+05
75%	6.450000e+05
max	7.700000e+06

Name: Sale Price, dtype: float64

```
In [5]: df.info()
```

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 21613 entries, 0 to 21612  
Data columns (total 21 columns):  
# Column Non-Null Count Dtype  
--- -  
0 ID 21613 non-null int64  
1 Date House was Sold 21613 non-null object  
2 Sale Price 21609 non-null float64  
3 No of Bedrooms 21613 non-null int64  
4 No of Bathrooms 21609 non-null float64  
5 Flat Area (in Sqft) 21604 non-null float64  
6 Lot Area (in Sqft) 21604 non-null float64  
7 No of Floors 21613 non-null float64  
8 Waterfront View 21613 non-null object  
9 No of Times Visited 21613 non-null object  
10 Condition of the House 21613 non-null object  
11 Overall Grade 21613 non-null int64  
12 Area of the House from Basement (in Sqft) 21610 non-null float64  
13 Basement Area (in Sqft) 21613 non-null int64  
14 Age of House (in Years) 21613 non-null int64  
15 Renovated Year 21613 non-null int64  
16 Zipcode 21612 non-null float64  
17 Latitude 21612 non-null float64  
18 Longitude 21612 non-null float64  
19 Living Area after Renovation (in Sqft) 21612 non-null float64  
20 Lot Area after Renovation (in Sqft) 21613 non-null int64  
dtypes: float64(10), int64(7), object(4)  
memory usage: 3.5+ MB

```
In [6]: df.describe()
```

Out[6]:

	ID	Sale Price	No of Bedrooms	No of Bathrooms	Flat Area (in Sqft)	Lot Area (in Sqft)	No of Floors	Overall Grade	Area of the House from Basement (in Sqft)	E Area
count	2.161300e+04	2.160900e+04	21613.000000	21609.000000	21604.000000	2.160400e+04	21613.000000	21613.000000	21610.000000	2161
mean	4.580302e+09	5.401984e+05	3.370842	2.114732	2079.931772	1.510776e+04	1.494309	7.623467	1788.344193	29
std	2.876566e+09	3.673890e+05	0.930062	0.770138	918.487597	4.142827e+04	0.539989	1.105439	827.982604	44
min	1.000102e+06	7.500000e+04	0.000000	0.000000	290.000000	5.200000e+02	1.000000	1.000000	290.000000	
25%	2.123049e+09	3.219500e+05	3.000000	1.750000	1429.250000	5.040000e+03	1.000000	7.000000	1190.000000	
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.617500e+03	1.500000	7.000000	1560.000000	
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068825e+04	2.000000	8.000000	2210.000000	56
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	10.000000	9410.000000	482

```
In [7]: zip_cond=df.groupby(['Condition of the House'])['Sale Price'].mean()
```

```
In [8]: zip_cond
```

```
Out[8]: Condition of the House
Bad      334431.666667
Excellent 612577.742504
Fair      542130.611206
Good      521277.510567
Okay     327316.215116
Name: Sale Price, dtype: float64
```

```
In [9]: l=[]
for i in zip_cond:
    l.append(i)
```

```
In [10]: l
```

```
Out[10]: [334431.6666666667,
612577.7425044092,
542130.6112061591,
521277.51056710107,
327316.2151162791]
```

```
In [11]: labels=df['Condition of the House'].unique()
```


```
In [12]: labels
```

```
Out[12]: array(['Fair', 'Excellent', 'Good', 'Bad', 'Okay'], dtype=object)
```


```
In [13]: import numpy as np
import plotly.express as px
# This dataframe has 244 lines, but 4 distinct values for `day`
fig = px.pie(df, values=l, names=labels, title='mean sale price as per condition of house', color_discrete_sequence=px.colors.qualitative.M3)
fig.show()
```



mean sale price as per condition of house



```
In [14]: #OUTLIERS IN GRAPHICAL VIEW
fig = px.scatter(df, x='ID', y="Sale Price")
fig.show()
```



```
In [15]: #IDENTIFY OUTLIERS IN MATHEMATICAL FORM with inter quartile method
```

```
q3=df['Sale Price'].quantile(.75)
q1=df['Sale Price'].quantile(.25)
```

```
In [16]: iqr=q3-q1
```

```
In [17]: iqr
```

```
Out[17]: 323050.0
```

```
In [18]: lower_limit=q1-1.5*iqr
```

```
In [19]: lower_limit
```

```
Out[19]: -162625.0
```

```
In [20]: upper_limit=q1+1.5*iqr
```


```
In [21]: #perform imputing method to treat with outliers(independent variables)
```

```
def limit_imputer(value):
    if value>upper_limit:
        return upper_limit
    if value<lower_limit:
        return lower_limit
    else:
        return value
df['Sale Price']=df['Sale Price'].apply(limit_imputer)
```

```
In [22]: df['Sale Price'].describe()
```

```
Out[22]: count      21609.000000
mean       485048.221667
std       197301.596234
min         75000.000000
25%       321950.000000
50%       450000.000000
75%       645000.000000
max       806525.000000
Name: Sale Price, dtype: float64
```

```
In [23]: fig = px.scatter(df, x='ID', y="Sale Price",trendline='ols')
fig.show()
```



```
In [24]: #treat with missing values(target variables)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0    ID                                              21613 non-null  int64
1    Date House was Sold                           21613 non-null  object
2    Sale Price                                    21609 non-null  float64
3    No of Bedrooms                               21613 non-null  int64
4    No of Bathrooms                              21609 non-null  float64
5    Flat Area (in Sqft)                          21604 non-null  float64
6    Lot Area (in Sqft)                           21604 non-null  float64
7    No of Floors                                 21613 non-null  float64
8    Waterfront View                              21613 non-null  object
9    No of Times Visited                          21613 non-null  object
10   Condition of the House                      21613 non-null  object
11   Overall Grade                               21613 non-null  int64
12   Area of the House from Basement (in Sqft)   21610 non-null  float64
13   Basement Area (in Sqft)                    21613 non-null  int64
14   Age of House (in Years)                    21613 non-null  int64
15   Renovated Year                             21613 non-null  int64
16   Zipcode                                    21612 non-null  float64
17   Latitude                                   21612 non-null  float64
18   Longitude                                  21612 non-null  float64
19   Living Area after Renovation (in Sqft)      21612 non-null  float64
20   Lot Area after Renovation (in Sqft)         21613 non-null  int64
dtypes: float64(10), int64(7), object(4)
memory usage: 3.5+ MB
```

```
In [25]: df.dropna(inplace=True,axis=0,subset=['Sale Price'])
```

```
In [26]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21609 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0    ID                                21609 non-null  int64
1    Date House was Sold              21609 non-null  object
2    Sale Price                       21609 non-null  float64
3    No of Bedrooms                   21609 non-null  int64
4    No of Bathrooms                 21605 non-null  float64
5    Flat Area (in Sqft)             21600 non-null  float64
6    Lot Area (in Sqft)              21600 non-null  float64
7    No of Floors                     21609 non-null  float64
8    Waterfront View                  21609 non-null  object
9    No of Times Visited              21609 non-null  object
10   Condition of the House           21609 non-null  object
11   Overall Grade                    21609 non-null  int64
12   Area of the House from Basement (in Sqft) 21606 non-null  float64
13   Basement Area (in Sqft)         21609 non-null  int64
14   Age of House (in Years)         21609 non-null  int64
15   Renovated Year                  21609 non-null  int64
16   Zipcode                         21608 non-null  float64
17   Latitude                        21608 non-null  float64
18   Longitude                       21608 non-null  float64
19   Living Area after Renovation (in Sqft) 21608 non-null  float64
20   Lot Area after Renovation (in Sqft) 21609 non-null  int64
dtypes: float64(10), int64(7), object(4)
memory usage: 3.6+ MB
```

```
In [30]: #treat mising values of (independent variables)
num_col=['No of Bathrooms','Flat Area (in Sqft)','Lot Area (in Sqft)','Area of the House from Basement (in Sqft)']
```

```
In [31]: from sklearn.impute import SimpleImputer
imputer=SimpleImputer(missing_values=np.nan,strategy='median')
df[num_col]=imputer.fit_transform(df[num_col])
```

```
In [32]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21609 entries, 0 to 21612
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0    ID                                21609 non-null  int64
1    Date House was Sold              21609 non-null  object
2    Sale Price                       21609 non-null  float64
3    No of Bedrooms                   21609 non-null  int64
4    No of Bathrooms                 21609 non-null  float64
5    Flat Area (in Sqft)             21609 non-null  float64
6    Lot Area (in Sqft)              21609 non-null  float64
7    No of Floors                     21609 non-null  float64
8    Waterfront View                  21609 non-null  object
9    No of Times Visited              21609 non-null  object
10   Condition of the House           21609 non-null  object
11   Overall Grade                    21609 non-null  int64
12   Area of the House from Basement (in Sqft) 21609 non-null  float64
13   Basement Area (in Sqft)         21609 non-null  int64
14   Age of House (in Years)         21609 non-null  int64
15   Renovated Year                  21609 non-null  int64
16   Zipcode                         21608 non-null  float64
17   Latitude                        21609 non-null  float64
18   Longitude                       21609 non-null  float64
19   Living Area after Renovation (in Sqft) 21609 non-null  float64
20   Lot Area after Renovation (in Sqft) 21609 non-null  int64
dtypes: float64(10), int64(7), object(4)
memory usage: 3.6+ MB
```

```
In [36]: #IMP: ZIPCODE CANT BE TREATED AS CONTINUOUS VARIABLE ,IT SHOULD BE TREATED AS CATEGORICAL VARIABLE SO,THIS WAY
imputer=SimpleImputer(missing_values=np.nan,strategy='most_frequent')
df['Zipcode']=imputer.fit_transform(df['Zipcode'])
#BUT HERE IS AN IMP ERROR TO UNDERSTAND
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[36], line 3
      1 #IMP: ZIPCODE CANT BE TREATED AS CONTINUOUS VARIABLE ,IT SHOULD BE TREATED AS CATEGORICAL VARIABLE SO,
      2 imputer=SimpleImputer(missing_values=np.nan,strategy='most_frequent')
----> 3 df['Zipcode']=imputer.fit_transform(df['Zipcode'])

File ~\anaconda3\lib\site-packages\sklearn\utils\_set_output.py:142, in _wrap_method_output.<locals>.wrapped(self, X, *args, **kwargs)
    140 @wraps(f)
    141 def wrapped(self, X, *args, **kwargs):
--> 142     data_to_wrap = f(self, X, *args, **kwargs)
    143     if isinstance(data_to_wrap, tuple):
    144         # only wrap the first output for cross decomposition
    145         return (
    146             _wrap_data_with_container(method, data_to_wrap[0], X, self),
```

```

147         *data_to_wrap[1:],
148     )

File ~\anaconda3\lib\site-packages\sklearn\base.py:859, in TransformerMixin.fit_transform(self, X, y, **fit_params)
855 # non-optimized default implementation; override when a better
856 # method is possible for a given clustering algorithm
857 if y is None:
858     # fit method of arity 1 (unsupervised transformation)
--> 859     return self.fit(X, **fit_params).transform(X)
860 else:
861     # fit method of arity 2 (supervised transformation)
862     return self.fit(X, y, **fit_params).transform(X)

File ~\anaconda3\lib\site-packages\sklearn\impute\_base.py:390, in SimpleImputer.fit(self, X, y)
381 if self.verbose != "deprecated":
382     warnings.warn(
383         "The 'verbose' parameter was deprecated in version "
384         "1.1 and will be removed in 1.3. A warning will "
385         (...)
386         FutureWarning,
387     )
--> 390 X = self._validate_input(X, in_fit=True)
392 # default fill_value is 0 for numerical input and "missing_value"
393 # otherwise
394 if self.fill_value is None:

File ~\anaconda3\lib\site-packages\sklearn\impute\_base.py:344, in SimpleImputer._validate_input(self, X, in_fit)
342         raise new_ve from None
343     else:
--> 344         raise ve
346 if in_fit:
347     # Use the dtype seen in `fit` for non-`fit` conversion
348     self._fit_dtype = X.dtype

File ~\anaconda3\lib\site-packages\sklearn\impute\_base.py:327, in SimpleImputer._validate_input(self, X, in_fit)
324     force_all_finite = True
326 try:
--> 327     X = self._validate_data(
328         X,
329         reset=in_fit,
330         accept_sparse="csc",
331         dtype=dtype,
332         force_all_finite=force_all_finite,
333         copy=self.copy,
334     )
335 except ValueError as ve:
336     if "could not convert" in str(ve):

File ~\anaconda3\lib\site-packages\sklearn\base.py:546, in BaseEstimator._validate_data(self, X, y, reset, validate_separately, **check_params)
544     raise ValueError("Validation should be done on X, y or both.")
545 elif not no_val_X and no_val_y:
--> 546     X = check_array(X, input_name="X", **check_params)
547     out = X
548 elif no_val_X and not no_val_y:

File ~\anaconda3\lib\site-packages\sklearn\utils\validation.py:902, in check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, estimator, input_name)
900     # If input is 1D raise error
901     if array.ndim == 1:
--> 902         raise ValueError(
903             "Expected 2D array, got 1D array instead:\narray={}. \n"
904             "Reshape your data either using array.reshape(-1, 1) if "
905             "your data has a single feature or array.reshape(1, -1) "
906             "if it contains a single sample.".format(array)
907         )
909 if dtype_numeric and array.dtype.kind in "USV":
910     raise ValueError(
911         "dtype='numeric' is not compatible with arrays of bytes/strings."
912         "Convert your data to numeric values explicitly instead."
913     )

ValueError: Expected 2D array, got 1D array instead:
array=[98178. 98125. 98028. ... 98144. 98027. 98144.].
Reshape your data either using array.reshape(-1, 1) if your data has a single feature or array.reshape(1, -1) if it contains a single sample.

```

In [37]: df['Zipcode'].shape

Out[37]: (21609,)

In [49]: #we need to convert our array in 2d  
column=df['Zipcode'].values.reshape(-1,1)

```

In [50]: column.shape
Out[50]: (21609, 1)

In [59]: #as our zipcode is in 2d,now we can do fit transform again for zipcode column
column=df['Zipcode'].values.reshape(-1,1)
imputer=SimpleImputer(missing_values=np.nan,strategy='most_frequent')
df['Zipcode']=imputer.fit_transform(column)

In [60]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21609 entries, 0 to 21612
Data columns (total 22 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   ID                                    21609 non-null  int64
 1   Date House was Sold                  21609 non-null  object
 2   Sale Price                           21609 non-null  float64
 3   No of Bedrooms                       21609 non-null  int64
 4   No of Bathrooms                     21609 non-null  float64
 5   Flat Area (in Sqft)                  21609 non-null  float64
 6   Lot Area (in Sqft)                   21609 non-null  float64
 7   No of Floors                         21609 non-null  float64
 8   Waterfront View                      21609 non-null  object
 9   No of Times Visited                  21609 non-null  object
10   Condition of the House               21609 non-null  object
11   Overall Grade                        21609 non-null  int64
12   Area of the House from Basement (in Sqft) 21609 non-null  float64
13   Basement Area (in Sqft)              21609 non-null  int64
14   Age of House (in Years)              21609 non-null  int64
15   Renovated Year                       21609 non-null  int64
16   Zipcode                              21609 non-null  float64
17   Latitude                             21609 non-null  float64
18   Longitude                            21609 non-null  float64
19   Living Area after Renovation (in Sqft) 21609 non-null  float64
20   Lot Area after Renovation (in Sqft)    21609 non-null  int64
21   zipcode                              21609 non-null  float64
dtypes: float64(11), int64(7), object(4)
memory usage: 3.8+ MB


In [61]: #variable transformon
#we cannot have direct relation with sale price and zip code so we will but we can check sale price zipcode wis
df['Zipcode']=df['Zipcode'].astype(object)
df.dtypes

Out[61]:
ID                                int64
Date House was Sold               object
Sale Price                        float64
No of Bedrooms                    int64
No of Bathrooms                   float64
Flat Area (in Sqft)               float64
Lot Area (in Sqft)                float64
No of Floors                      float64
Waterfront View                   object
No of Times Visited               object
Condition of the House             object
Overall Grade                     int64
Area of the House from Basement (in Sqft) float64
Basement Area (in Sqft)           int64
Age of House (in Years)           int64
Renovated Year                    int64
Zipcode                           object
Latitude                          float64
Longitude                         float64
Living Area after Renovation (in Sqft) float64
Lot Area after Renovation (in Sqft) int64
zipcode                           float64
dtype: object

In [70]: import plotly.express as px
fig = px.bar(df, x = "Zipcode", y = "Sale Price",title='variation of sale price as per zipcode')
fig.show()

```

## variation of sale price as per zipcode



```
In [73]: #variable transformation over 'no of times visited' col as it does not represent data properly
df['No of Times Visited'].unique()
```

```
Out[73]: array(['None', 'Thrice', 'Four', 'Twice', 'Once'], dtype=object)
```

```
In [74]: mapping={'None':'0', 'Thrice':'3', 'Four':'4', 'Twice':'2', 'Once':'1'}
```

```
In [75]: df['No of Times Visited']=df['No of Times Visited'].map(mapping)
```

```
In [76]: df['No of Times Visited'].unique()
```

```
Out[76]: array(['0', '3', '4', '2', '1'], dtype=object)
```

```
In [77]: #variable transformation used to create new variable by combining or transforming 2 variables
#by looking at 'renovation year' we can have que when renovation was done and how it can impact sale price
#create 2 variables 'ever renovated' and 'years since renovataion'
df['ever renovated']=np.where(df['Renovated Year']==0,'No','Yes')
```

```
In [80]: #for 2nd variable fetch year from 'date of sold col'
df['purchase year']=pd.DatetimeIndex(df['Date House was Sold']).year
```

```
In [82]: df['purchase year'].head(5)
```

```
Out[82]:
0    2017
1    2017
2    2016
3    2017
4    2016
Name: purchase year, dtype: int64
```

```
In [91]: df['years since renovataion']=np.where(df['ever renovated']=='Yes',abs(df['purchase year']-df['Renovated Year'])
```

```
In [92]: df['years since renovataion'].tail(5)
```

```
Out[92]:
21608    0
21609    0
21610    0
21611    0
21612    0
Name: years since renovataion, dtype: int64
```

```
In [93]: df.head()
```



Out[93]:

	ID	Date House was Sold	Sale Price	No of Bedrooms	No of Bathrooms	Flat Area (in Sqft)	Lot Area (in Sqft)	No of Floors	Waterfront View	No of Times Visited	...	Renovated Year	Zipcode	Latitude	Longitude
0	7129300520	14 October 2017	221900.0	3	1.00	1180.0	5650.0	1.0	No	0	...	0	98178.0	47.5112	
1	6414100192	14 December 2017	538000.0	3	2.25	2570.0	7242.0	2.0	No	0	...	1991	98125.0	47.7210	
2	5631500400	15 February 2016	180000.0	2	1.00	770.0	10000.0	1.0	No	0	...	0	98028.0	47.7379	
3	2487200875	14 December 2017	604000.0	4	3.00	1960.0	5000.0	1.0	No	0	...	0	98136.0	47.5208	
4	1954400510	15 February 2016	510000.0	3	2.00	1680.0	8080.0	1.0	No	0	...	0	98074.0	47.6168	

5 rows × 25 columns

In [94]:

```
#AS WE SEE WE GOT OUR COLOUMN NOW WE DONT REQUIRE USELESS COLUMNS,WE WILL DELETE THOSE TO KEEP OUR DATA TIDY
df.drop(columns=['Date House was Sold','Renovated Year','purchase year'],inplace=True)
```

In [95]:

```
df.head()
```

Out[95]:

	ID	Sale Price	No of Bedrooms	No of Bathrooms	Flat Area (in Sqft)	Lot Area (in Sqft)	No of Floors	Waterfront View	No of Times Visited	Condition of the House	...	Basement Area (in Sqft)	Age of House (in Years)	Zipcode	Latitude	Longitude
0	7129300520	221900.0	3	1.00	1180.0	5650.0	1.0	No	0	Fair	...	0	63	98178.0	47.5112	
1	6414100192	538000.0	3	2.25	2570.0	7242.0	2.0	No	0	Fair	...	400	67	98125.0	47.7210	
2	5631500400	180000.0	2	1.00	770.0	10000.0	1.0	No	0	Fair	...	0	85	98028.0	47.7379	
3	2487200875	604000.0	4	3.00	1960.0	5000.0	1.0	No	0	Excellent	...	910	53	98136.0	47.5208	
4	1954400510	510000.0	3	2.00	1680.0	8080.0	1.0	No	0	Fair	...	0	31	98074.0	47.6168	

5 rows × 22 columns

In [96]:

```
#correlation calculation
df['Sale Price'].corr(df['Flat Area (in Sqft)'])
```

Out[96]:

0.6492472259786739

In [97]:

```
#correlation between independent variables
df.drop(columns=['ID']).corr()
```

C:\Users\Janhavi\AppData\Local\Temp\ipykernel\_8236\3805217886.py:2: FutureWarning:  
The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

Out[97]:

	Sale Price	No of Bedrooms	No of Bathrooms	Flat Area (in Sqft)	Lot Area (in Sqft)	No of Floors	Overall Grade	Area of the House from Basement (in Sqft)	Basement Area (in Sqft)	Age of House (in Years)	Latitude	Longitude
Sale Price	1.000000	0.333635	0.509890	0.649247	0.101291	0.303304	0.664972	0.568595	0.283494	-0.066771	0.453716	0.063154
No of Bedrooms	0.333635	1.000000	0.515813	0.576628	0.031692	0.175536	0.349223	0.477549	0.303294	-0.154113	-0.008708	0.129569
No of Bathrooms	0.509890	0.515813	1.000000	0.754568	0.087732	0.500776	0.635638	0.685088	0.283798	-0.505954	0.024570	0.223171
Flat Area (in Sqft)	0.649247	0.576628	0.754568	1.000000	0.172721	0.354142	0.705725	0.876226	0.435142	-0.318146	0.052538	0.240091
Lot Area (in Sqft)	0.101291	0.031692	0.087732	0.172721	1.000000	-0.005162	0.102314	0.183492	0.015252	-0.053119	-0.085719	0.229449
No of Floors	0.303304	0.175536	0.500776	0.354142	-0.005162	1.000000	0.461368	0.524031	-0.245572	-0.489244	0.049692	0.125620
Overall Grade	0.664972	0.349223	0.635638	0.705725	0.102314	0.461368	1.000000	0.705153	0.145232	-0.456711	0.111226	0.201736
Area of the House from Basement (in Sqft)	0.568595	0.477549	0.685088	0.876226	0.183492	0.524031	0.705153	1.000000	-0.051825	-0.423848	-0.000819	0.343793
Basement Area (in Sqft)	0.283494	0.303294	0.283798	0.435142	0.015252	-0.245572	0.145232	-0.051825	1.000000	0.133072	0.110451	-0.144822
Age of House (in Years)	-0.066771	-0.154113	-0.505954	-0.318146	-0.053119	-0.489244	-0.456711	-0.423848	0.133072	1.000000	0.148083	-0.409515
Latitude	0.453716	-0.008708	0.024570	0.052538	-0.085719	0.049692	0.111226	-0.000819	0.110451	0.148083	1.000000	-0.135551
Longitude	0.063154	0.129569	0.223171	0.240091	0.229449	0.125620	0.201736	0.343793	-0.144822	-0.409515	-0.135551	1.000000
Living Area after Renovation (in Sqft)	0.600540	0.391771	0.568568	0.756185	0.144507	0.280106	0.681362	0.731996	0.200302	-0.326307	0.048836	0.334507
Lot Area after Renovation (in Sqft)	0.092041	0.029264	0.087226	0.183223	0.718527	-0.011204	0.107581	0.194106	0.017263	-0.071016	-0.086420	0.254449
zipcode	-0.033636	-0.152760	-0.203951	-0.199380	-0.129551	-0.059222	-0.185844	-0.261124	0.074933	0.346928	0.267022	-0.564107
years since renovataion	0.054093	-0.007198	0.003551	0.023503	0.013835	-0.000901	-0.024388	0.010491	0.029158	0.203375	0.019739	-0.055000

In [98]: `#how to identify categorical variable: variable with datatype object`  
`df.info()`

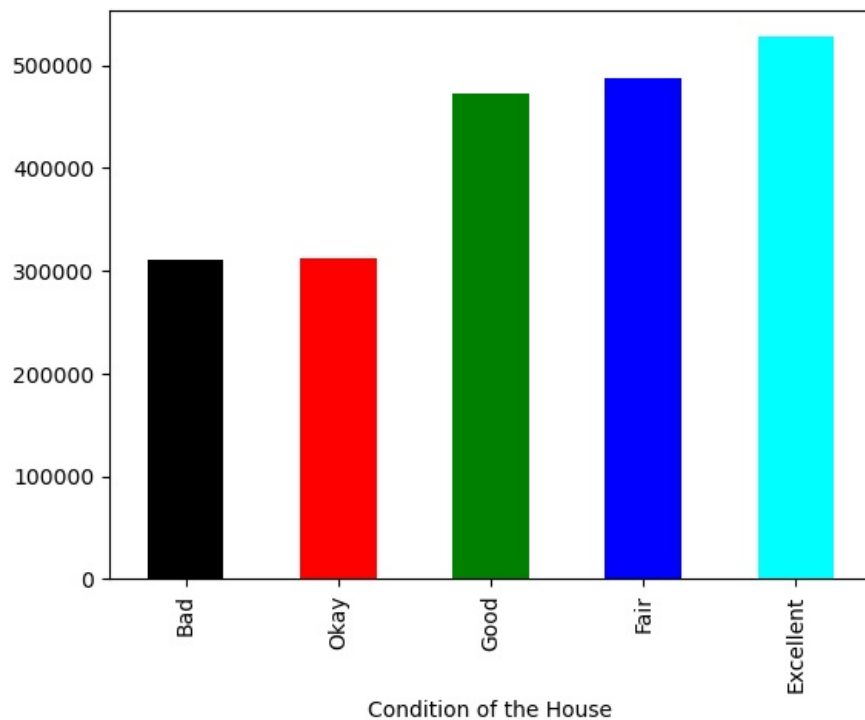
```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21609 entries, 0 to 21612
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                    21609 non-null  int64
1   Sale Price                           21609 non-null  float64
2   No of Bedrooms                       21609 non-null  int64
3   No of Bathrooms                     21609 non-null  float64
4   Flat Area (in Sqft)                 21609 non-null  float64
5   Lot Area (in Sqft)                 21609 non-null  float64
6   No of Floors                        21609 non-null  float64
7   Waterfront View                     21609 non-null  object
8   No of Times Visited                 21609 non-null  object
9   Condition of the House              21609 non-null  object
10  Overall Grade                       21609 non-null  int64
11  Area of the House from Basement (in Sqft) 21609 non-null  float64
12  Basement Area (in Sqft)             21609 non-null  int64
13  Age of House (in Years)             21609 non-null  int64
14  Zipcode                             21609 non-null  object
15  Latitude                            21609 non-null  float64
16  Longitude                           21609 non-null  float64
17  Living Area after Renovation (in Sqft) 21609 non-null  float64
18  Lot Area after Renovation (in Sqft)  21609 non-null  int64
19  zipcode                             21609 non-null  float64
20  ever renovated                      21609 non-null  object
21  years since renovataion              21609 non-null  int64
dtypes: float64(10), int64(7), object(5)
memory usage: 3.8+ MB
```

In [99]: `df['Condition of the House'].value_counts()`

```
Out[99]: Fair      14028
Good       5678
Excellent  1701
Okay      172
Bad        30
Name: Condition of the House, dtype: int64
```

```
In [114]: #relationship of independent with dependent variables
df.groupby('Condition of the House')['Sale Price'].mean().sort_values().plot(kind='bar',color=['black', 'red',
```

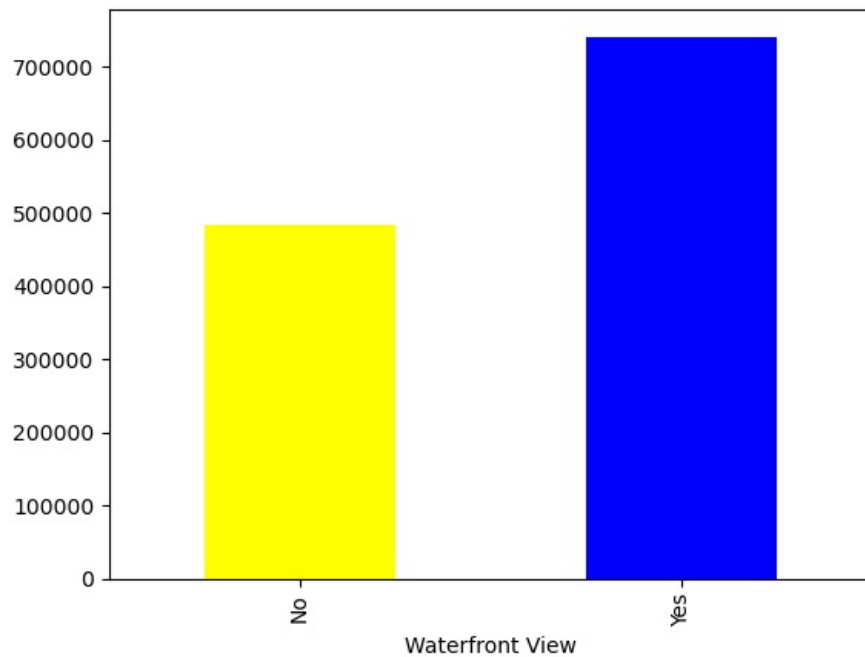
```
Out[114]: <Axes: xlabel='Condition of the House'>
```



```
In [ ]:
```

```
In [117]: df.groupby('Waterfront View')['Sale Price'].mean().sort_values().plot(kind='bar',color=['yellow', 'blue'])
```

```
Out[117]: <Axes: xlabel='Waterfront View'>
```



```
In [118]: df.groupby('Zipcode')['Sale Price'].mean().sort_values().plot(kind='bar')
```

```
Out[118]: <Axes: xlabel='Zipcode'>
```

